

Logistic Regression

CS 584: Data Mining

Prof. Sanmay Das (Adapted from notes of Prof. Malik Magdon-Ismael)

Recap + Setting

- Linear regression: Find a weight vector \mathbf{w} that

minimizes
$$E_{\text{in}}(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n (\mathbf{w} \cdot \mathbf{x}_i - f(\mathbf{x}_i))^2$$

- $h(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x}$

- Classification problem: Will someone have a heart attack over the next year?

age	62 years
gender	male
blood sugar	120 mg/dL40,000
HDL	50
LDL	120
Mass	190 lbs
Height	5' 10''
...	...

- Unknown target function $f : \mathcal{X} \rightarrow \mathcal{Y}$
 - ▶ Classification is when \mathcal{Y} is categorical (e.g. binary)
- Training data $\mathcal{D} : (\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)$ where $y_i = f(\mathbf{x}_i)$ (possibly noisy).
- Want to learn h “close to” f .
- Two central questions:
 - ▶ How do we learn h ?
 - ★ Key algorithmic question!
 - ▶ What can we say about how close h is to f ?
 - ★ Why is this hard?

- Logistic regression: Predict probability of a heart attack. $y \in [0,1]$

• $h(\mathbf{x}) = \sigma(\mathbf{w} \cdot \mathbf{x})$ where $\sigma(z) = \frac{1}{1 + e^{-z}}$




Figure 4.19. Plot of sigmoid (logistic) function, $\sigma(z)$.

Understanding the Setting

- The data is still binary. $y \in \{-1, 1\}$
- We never see the *probability* $f(\mathbf{x}) = \Pr[y = 1 \mid \mathbf{x}]$ of someone having a heart attack
 - Just whether they *did* or not
 - Ways to think about this:
 - Each individual has a risk state or subjective probability, and whether the outcome happens is based on a biased coin flip
 - Each individual has a true state $\in \{-1, +1\}$ and then the data is generated from a *noisy* target function $\Pr[y = 1 \mid \mathbf{x}] = f(\mathbf{x})$

Finding a Good Hypothesis

- Our hypothesis is basically just a vector \mathbf{w} — the logistic transformation is pre-specified.
- Two notes:
 - I will use the convention of just using \mathbf{w} rather than b and \mathbf{w} separately. The simplest way to think about this is just of augmenting the \mathbf{x} vector with a first element that always has the value 1 for any i .
 - We are using labels -1 and +1. The book uses 0, 1 for logistic regression. This leads to minor differences in formulas.
- What does it mean to find a “good” \mathbf{w} in this setting?
- Two ways to see it:
 - Minimizing an error measure (conventional ML view)
 - Maximizing the probability of seeing the actual y values we saw

Probabilistic Interpretation

- $\mathcal{L}(\mathbf{w}) = \prod_{i=1}^n \Pr(y_i | \mathbf{x}_i, \mathbf{w})$
- We want to pick the \mathbf{w} that maximizes this. Let's walk through this:

- Same as maximizing: $\ln \mathcal{L}(\mathbf{w}) = \sum_{i=1}^n \ln \Pr(y_i | \mathbf{x}_i, \mathbf{w})$

- Or, minimizing: $-\ln \mathcal{L}(\mathbf{w})$.

- Which is the same as minimizing: $-\frac{1}{n} \sum_{i=1}^n \ln \Pr(y_i | \mathbf{x}_i, \mathbf{w})$.

- I can write this minimization as: $\min \frac{1}{n} \sum_{i=1}^n \ln \frac{1}{\Pr(y_i | \mathbf{x}_i, \mathbf{w})}$

Minimizing the log likelihood

- When $y = +1$, $\Pr(y | \mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w} \cdot \mathbf{x}) = \sigma(y_i \mathbf{w} \cdot \mathbf{x})$
- When $y = -1$, $\Pr(y | \mathbf{x}, \mathbf{w}) = 1 - \sigma(\mathbf{w} \cdot \mathbf{x}) = \sigma(-\mathbf{w} \cdot \mathbf{x}) = \sigma(y_i \mathbf{w} \cdot \mathbf{x})$

- So we want to minimize $\frac{1}{n} \sum_{i=1}^n \ln \frac{1}{\sigma(y_i \mathbf{w} \cdot \mathbf{x}_i)} = \frac{1}{n} \sum_{i=1}^n \ln(1 + e^{-y_i \mathbf{w} \cdot \mathbf{x}_i})$

- This is also known as the **cross entropy error**

- Exercises:

- Error if you “predict” 0.9 on a +1 example? -1 example? $\approx 0.105, 2.303$
- What if you “predict” 0.7 on a +1 example? -1 example? $\approx 0.357, 1.204$

Low when y_i has
the same sign as
 $\mathbf{w} \cdot \mathbf{x}_i$ and is large

Minimizing Cross-Entropy Error

- Turns out to be a convex function, therefore has a unique minimum.
- Can use conventional optimization methods (e.g. Newton-Raphson).
- Later today we'll talk about how to do this using *gradient descent*, a key technique in modern ML and DM.